

```

import torch
import torch.nn.functional as F          # Para usar linear.
import torchvision as tv

class RBM( torch.nn.Module):
    def __init__( _, vsize, hsize, CD_k=1):
        super().__init__()
        _ .W = torch.nn.Parameter( torch.randn( hsize, vsize)*1e-2)
        _ .bv = torch.nn.Parameter( torch.randn( vsize)*1e-2) # Bias V.
        _ .bh = torch.nn.Parameter( torch.randn( hsize)*1e-2) # Bias H.
        _ .k = CD_k          # Divergencia Contrastiva.

    def sample_h( _, v):
        prob_h = torch.sigmoid( F.linear( v, _ .W, _ .bh))
        samp_h = torch.bernoulli( prob_h)
        return prob_h, samp_h

    def sample_v( _, h):
        prob_v = torch.sigmoid( F.linear( h, _ .W.t(), _ .bv))
        samp_v = torch.bernoulli( prob_v)
        return prob_v, samp_v

    def forward( _, v):
        vs = v
        for i in range( _ .k):
            hp, hs = _ .sample_h( vs)
            vp, vs = _ .sample_v( hs)
        return v, vs

    def free_energy( _, v):
        v_bv = v.mv( _ .bv)          # Multiplica matriz por vector.
        hlin = torch.clamp( F.linear( v, _ .W, _ .bh), -80, 80)
        slog = hlin.exp().add(1).log().sum(1)
        return ( -slog - v_bv).mean()

if __name__ == '__main__':
    # Para poder importar RBM.
    T = 20          # Cant de epocas.
    B = 64          # Mini-lotes.
                    # trn/tst-data/load como antes.
    N = 28*28       # Cant de entradas.
    M = 64          # Cant de features.
    C = 10          # Cant de clases de salida.

```

```

rbmfd = RBM( N, M)          # Feature Detector
optim = torch.optim.SGD( rbmfd.parameters(), 0.1)
for t in range(T):
    E = 0
    for images, labels in trn_load:
        optim.zero_grad()
        data = images.view( -1, N)
        v0, vk = rbmfd( data)
        loss = rbmfd.free_energy(v0) - rbmfd.free_energy(vk)
        loss.backward()
        optim.step()
        E += loss.item()
    print( t, E)

lincl = torch.nn.Linear( M, C)          # Linear Classifier
optim = torch.optim.SGD( lincl.parameters(), 0.1)
costf = torch.nn.CrossEntropyLoss()
for t in range(T):
    E = 0
    for images, labels in trn_load:
        optim.zero_grad()
        v = images.view( -1, N)
        hp, hs = rbmfd.sample_h( v)
        cp = lincl( hp)
        error = costf( cp, labels)
        error.backward()
        optim.step()
        E += error.item()
    print( t, E)

    right = 0
    total = 0
    with torch.no_grad():
        for images, labels in tst_load:
            v = images.view( -1, N)
            hp, hs = rbmfd.sample_h( v)
            cp = lincl( hp)
            right += (cp.argmax(dim=1)==labels).sum().item()
            total += len(labels)
    print("Accuracy:", right/total)

```

Ver: <http://deeplearning.net/tutorial/rbm.html>